

Efficient Gaussian Processes for Model-based Online Planning

Hu Hanyang

Supervisor: Jonathan Scarlett

Department of Mathematics

National University of Singapore

Mathematics Capstone Project for Semester 2, AY2024/2025

April 15, 2025

Outline

- 1 Introduction
- 2 Backgrounds
- 3 Methods
- 4 Results
- 5 Potential Extensions
- 6 Conclusion

Introduction

In **model-free reinforcement learning**, 1 million time-steps is common for training, which might be infeasible for real-world applications.

Model-based reinforcement learning (MBRL), particularly online planning, may converge much earlier than 200k time-steps.¹

¹Tingwu Wang et al. *Benchmarking Model-Based Reinforcement Learning*. 2019.
URL: <https://arxiv.org/abs/1907.02057>.

Introduction

In **model-free reinforcement learning**, 1 million time-steps is common for training, which might be infeasible for real-world applications.

Model-based reinforcement learning (MBRL), particularly online planning, may converge much earlier than 200k time-steps.¹

Question: Can we make MBRL more **sample-efficient** by replacing the common MLP dynamics model with **Gaussian processes (GPs)**?

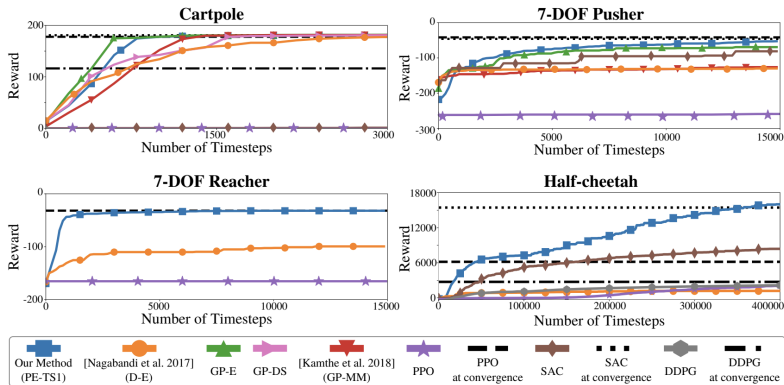
Challenges of GP dynamics for online planning:

- computational complexity (i.e., slow training and inference)
- curse of dimensionality (CoD)

¹Tingwu Wang et al. *Benchmarking Model-Based Reinforcement Learning*. 2019.
URL: <https://arxiv.org/abs/1907.02057>.

Existing Works on GP-based Planning

Results in the PE-TS² paper:




²Kurtland Chua et al. *Deep Reinforcement Learning in a Handful of Trials using Probabilistic Dynamics Models*. 2018.

Existing Works on GP-based Planning

Examples focusing on GP-based Planning:

- ① (Hewing et al., 2020) Propagating uncertainty by first-order approximations (similar to extended Kalman filters) and selecting inducing points dynamically.
- ② (Bosch et al., 2020) Using a neural network auto-encoder to alleviate the CoD, GP dynamics then plan in the latent space.



Note. These methods are typically tested only in simple environments.

Overarching Goal: Extend GP-based planning to more diverse domains while maintaining real-time performance and advantage over NN models.


Outline

- 1 Introduction
- 2 Backgrounds**
- 3 Methods
- 4 Results
- 5 Potential Extensions
- 6 Conclusion

Markov Decision Processes

A **Markov decision process (MDP)** is defined by $\mathcal{M} = (\mathcal{S}, \mathcal{A}, M, r, \rho_0)$

- state space \mathcal{S}
- action space \mathcal{A}
- transition probability distribution (dynamics) $s_{t+1} \sim M(\cdot | s_t, a_t)$
- reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$
- initial state distribution $s_0 \sim \rho_0(\cdot)$



Source: OpenAI Spinning Up (spinningup.openai.com).

Markov Decision Processes

Goal: Find a (deterministic) **policy** function $\pi : \mathcal{S} \rightarrow \mathcal{A}$ such that

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim P^{\pi}(\cdot)} [R(\tau)]$$

where $R(\cdot)$ denotes the **infinite-horizon discounted return**

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

with discount factor $\gamma \in (0, 1)$.

Markov Decision Processes

Goal: Find a (deterministic) **policy** function $\pi : \mathcal{S} \rightarrow \mathcal{A}$ such that

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim P^{\pi}(\cdot)} [R(\tau)]$$

where $R(\cdot)$ denotes the **infinite-horizon discounted return**

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

with discount factor $\gamma \in (0, 1)$.

Question: How to maximize this objective when we only have a set of sampled trajectories with finite length collected from a non-optimal policy?

Bellman Equation

Consider the **action-value function** $Q : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim P^\pi(\cdot)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \mid s_0 = s, a_0 = a \right]$$

which satisfies the **Bellman equation**

$$Q^\pi(s, a) = r(s, a) + \gamma \cdot \mathbb{E}_{s' \sim M(\cdot | s, a)} [Q^\pi(s', \pi(s'))].$$

Bellman Equation

Consider the **action-value function** $Q : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim P^\pi(\cdot)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \mid s_0 = s, a_0 = a \right]$$

which satisfies the **Bellman equation**

$$Q^\pi(s, a) = r(s, a) + \gamma \cdot \mathbb{E}_{s' \sim M(\cdot|s,a)} [Q^\pi(s', \pi(s'))].$$

The **optimal action-value function** Q^* should satisfy

$$Q^*(s, a) = r(s, a) + \gamma \cdot \mathbb{E}_{s' \sim M(\cdot|s,a)} [\max_{a \in \mathcal{A}} Q^*(s', a)]$$

and hence the optimal policy can be extracted by $\pi^*(s) = \max_a Q^*(s, a)$.

Deep Deterministic Policy Gradient

(1) Given a policy π , we can train a Q-network Q_θ by minimizing

$$L(\theta, \mathcal{D}) = \mathbb{E}_{(s_t, a_t, s_{t+1}, r_t) \sim \mathcal{D}} [\underbrace{(Q_\theta(s_t, a_t) - (r_t + \gamma Q_\theta(s_{t+1}, \pi(s_{t+1}))))^2}_{\text{temporal-difference (TD) target}}]$$

so that Q_θ approximates Q^π (referred to as **TD learning**).

(2) Find the policy network $\pi_\theta(s) \approx \arg \max_a Q^\pi(s, a)$ by maximizing

$$\mathbb{E}_{s \sim \mathcal{D}} [Q_{\theta_{\text{targ}}}(s, \pi_\theta(s))]$$

where θ_{targ} is a lagged target Q-network introduced to stabilize training.

The **deterministic policy gradient** theorem (Silver et al., 2014) states that this is approximately the same as the original objective of the MDP.

Intuition: RL is difficult, regression is easy.

With a learned dynamics model M_θ and reward model R_θ , the optimal action can be selected by **model predictive control (MPC)**

$$\pi_{\text{MPC}}(s_t) = \arg \max_{a_t} \max_{a_{t+1:t+H}} \mathbb{E} \left[\sum_{i=0}^H \gamma^i R_\theta(s_{t+i}, a_{t+i}) \right].$$

Model-based Online Planning

Algorithm 1 Model Predictive Control (MPC)

```
1: Input: Number of iterations  $J$ , population size  $N$ , number of elite samples  $K$ , roll-  
out horizon  $H$ , initial distribution parameters  $\mu^0, \Sigma^0$ , (learned) dynamics model  $M_\theta$ ,  
(learned) reward model  $R_\theta$ , current state  $s_t$ .  
2: for each iteration  $i = 1, 2, \dots, J$  do  
3:   Sample  $N$  action sequences of length  $H$  from  $\mathcal{N}(\mu^{j-1}, \Sigma^{j-1})$ .  
4:   for all  $N$  sequences  $\Gamma = (a_t, a_{t+1}, \dots, a_{t+H})$  do  
5:     for step  $j = 0, 1, \dots, H - 1$  do ▷ Estimate trajectory return  $\phi_\Gamma$   
6:       Update  $\phi_\Gamma = \phi_\Gamma + \gamma^t R_\theta(s_{t+j}, a_{t+j})$ . ▷ Initially setting  $\phi_\Gamma = 0$   
7:       Predict  $s_{t+j+1} \sim M_\theta(s_{t+j}, a_{t+j})$ .  
8:     end for  
9:   end for  
10:  Select the elite samples  $\{\Gamma_k^*\}_{k=1}^K$  corresponding to the top- $K$  returns  $\{\phi_{\Gamma_k^*}\}_{k=1}^K$ .  
11:  Update parameters  $\mu^j, \Sigma^j$  for the next iteration based on  $\{\Gamma_k^*\}_{k=1}^K$  and  $\{\phi_{\Gamma_k^*}\}_{k=1}^K$ .  
12: end for  
13: Output:  $(a_t^*, a_{t+1}^*, \dots, a_{t+H}^*) \sim \mathcal{N}(\mu^J, \Sigma^J)$ 
```

Time Complexity: $O(JNH)$ model inference steps. J is the number of iterations, N is the population size, and H is the planning horizon.

TD Learning for MPC (TD-MPC)

Intuition: Use TD learning to model global optimality and use MPC to refine local behaviors, requiring a shorter planning horizon H .

Modified MPC Objective:

$$\pi_{\text{TD-MPC}}(s_t) = \arg \max_{a_t} \max_{a_{t+1:t+H}} \mathbb{E} \left[\sum_{i=0}^{H-1} \gamma^i R_{\theta}(s_{t+i}, a_{t+i}) + \gamma^H Q_{\theta}(s_H, a_H) \right].$$

Note. All components of TD-MPC are implemented using deterministic neural networks.

Model-based Value Error of TD-MPC

Motivation: The model-based value error partially determines the performance of model-based online planning.

Theorem (modified and extended from Xiao et al., Theorem 1)

The model-based value error is bounded by

$$\left| V^\pi(s) - \hat{V}^\pi(s) \right| \leq \underbrace{K_{\mathcal{M}} \frac{\gamma - \gamma^{H+1}}{1 - \gamma} \epsilon_m}_{\text{dynamics gap}} + \underbrace{\frac{1 - \gamma^H}{1 - \gamma} \epsilon_r + \gamma^H \epsilon_q}_{\text{return estimation gap}}$$

and $K_{\mathcal{M}} \leq (L_R + 2\gamma V_{\max} L_M) \sqrt{1 + L_\pi^2}$.

Notations:

- ① **Dynamics Model Error** $\max_{s,a} W(M(\cdot|s, a), \hat{M}(\cdot|s, a)) \leq \epsilon_m$;
- ② **Reward Model Error** $\max_{s,a} |r(s, a) - \hat{r}(s, a)| \leq \epsilon_r$;
- ③ **Value Function Error** $\max_{s,a} |Q^\pi(s, a) - \hat{Q}^\pi(s, a)| \leq \epsilon_q$.

Outline

- 1 Introduction
- 2 Backgrounds
- 3 Methods**
- 4 Results
- 5 Potential Extensions
- 6 Conclusion

Key Design Choices:

- 1 Integration with TD-MPC.
- 2 GP-based Correction of MLPs.
- 3 Decoupled Training and Inference.
- 4 Integration with Deep Kernel Learning (DKL).



Note. We would refer to our method as GP-TD-MPC.

GP-based Correction

Instead of standalone GPs, we use an MLP model $f_\theta(\cdot)$ as the prior mean

$$\text{output}(\mathbf{x}^*) = f_\theta(\mathbf{x}^*) + \underbrace{\mathbf{k}_{X\mathbf{x}^*}^\top \hat{K}_{XX}^{-1}(\mathbf{y} - f_\theta(X))}_{\text{GP correction}}$$

and the GP training target becomes the residual $\mathbf{y} - f_\theta(X)$.



Legend: TD-MPC (black circles), GP-TD-MPC (residual target) (red squares), GP-TD-MPC (ground-truth target) (yellow diamonds), GP-TD-MPC (w/o MLP model) (purple triangles)

Computational Bottlenecks

Training a GP model involves maximizing the marginal log-likelihood (MLL) w.r.t. the kernel hyperparameters θ (e.g., the lengthscales)

$$\mathcal{L} = \log p(\mathbf{y} | X, \theta) \propto - \underbrace{(\mathbf{y} - m(X))^{\top} \hat{K}_{XX}^{-1} (\mathbf{y} - m(X))}_{\text{model fit}} - \underbrace{\log |\hat{K}_{XX}|}_{\text{complexity}}$$

which is computation-heavy.

Cached inference of the GP correction

$$\text{output}(\mathbf{x}^*) = f_{\theta}(\mathbf{x}^*) + \underbrace{\mathbf{k}_{X\mathbf{x}^*}^{\top} \hat{K}_{XX}^{-1} (\mathbf{y} - f_{\theta}(X))}_{\text{GP correction}}$$

takes $O(n)$ time for each independent GP (n is the size of the data).

Stochastic Variational Gaussian Process (SVGP)

Evidence Lower Bound (ELBO): Using a **variational distribution** $q(\mathbf{u}) \sim \mathcal{N}(\mathbf{m}_{\mathbf{u}}, S_{\mathbf{u}})$ w.r.t. m inducing points Z to approximate $p(\mathbf{f}, \mathbf{u}) = p(\mathbf{f} | \mathbf{u})p(\mathbf{u}) \approx p(\mathbf{f} | \mathbf{u})q(\mathbf{u})$ so that

$$\log p(\mathbf{y} | X, \theta) \geq \mathbb{E}_{q(\mathbf{u})p(\mathbf{f} | \mathbf{u})} [\log p(\mathbf{y} | \mathbf{f})] - D_{\text{KL}}(q(\mathbf{u}) \| p(\mathbf{u}))$$

where $p(\mathbf{y} | \mathbf{f}) = \prod_{i=1}^n p(\mathbf{y}_i | \mathbf{f}_i)$ usually factorizes over data instances.

Stochastic Variational Gaussian Process (SVGP)

Evidence Lower Bound (ELBO): Using a **variational distribution** $q(\mathbf{u}) \sim \mathcal{N}(\mathbf{m}_{\mathbf{u}}, S_{\mathbf{u}})$ w.r.t. m inducing points Z to approximate $p(\mathbf{f}, \mathbf{u}) = p(\mathbf{f} | \mathbf{u})p(\mathbf{u}) \approx p(\mathbf{f} | \mathbf{u})q(\mathbf{u})$ so that

$$\log p(\mathbf{y} | X, \theta) \geq \mathbb{E}_{q(\mathbf{u})p(\mathbf{f} | \mathbf{u})} [\log p(\mathbf{y} | \mathbf{f})] - D_{\text{KL}}(q(\mathbf{u}) \| p(\mathbf{u}))$$

where $p(\mathbf{y} | \mathbf{f}) = \prod_{i=1}^n p(\mathbf{y}_i | \mathbf{f}_i)$ usually factorizes over data instances.

The inference can also be approximated by

$$\begin{aligned}\mathbb{E}[f(\mathbf{x}^*) | X, \mathbf{y}, \mathbf{x}^*] &\approx \mathbf{k}_{Z\mathbf{x}^*}^{\top} K_{ZZ}^{-1} \mathbf{m}_{\mathbf{u}} \\ \text{Var}[f(\mathbf{x}^*) | X, \mathbf{y}, \mathbf{x}^*] &\approx k(\mathbf{x}^*, \mathbf{x}^*) - \mathbf{k}_{Z\mathbf{x}^*}^{\top} K_{ZZ}^{-1} (K_{ZZ} - S_{\mathbf{u}}) K_{ZZ}^{-1} \mathbf{k}_{Z\mathbf{x}^*}\end{aligned}$$


where the mean prediction only takes $O(m)$ time.

Note. We consider $m \equiv 0$ for simplicity.

Decoupled Training and Inference

Motivation: The SVGP training introduces $O(m^2)$ additional parameters, which we might want to avoid to reduce training time.

Training: Subsample a mini-batch of data for scalable training.³




³Shifan Zhao et al. *Efficient Two-Stage Gaussian Process Regression Via Automatic Kernel Search and Subsampling*. 2024.

Decoupled Training and Inference

Motivation: The SVGP training introduces $O(m^2)$ additional parameters, which we might want to avoid to reduce training time.

Training: Subsample a mini-batch of data for scalable training.³




Inference: Choose Z by **farthest point sampling** or **pivoted Cholesky**, then obtain the optimal variational distribution in $O(nm^2)$ time

$$\begin{aligned}\mathbf{c} &= K_{ZX}\Sigma_y^{-1}(\mathbf{y} - f_\theta(X)), & C &= K_{ZX}\Sigma_y^{-1}K_{XZ} \\ \mathbf{m}_u &= K_{ZZ}(K_{ZZ} + C)^{-1}\mathbf{c}, & S_u &= K_{ZZ}(K_{ZZ} + C)^{-1}K_{ZZ}.\end{aligned}$$

³Shifan Zhao et al. *Efficient Two-Stage Gaussian Process Regression Via Automatic Kernel Search and Subsampling*. 2024.

Deep Kernel Learning


In non-stationary and/or high-dimensional settings, we may introduce a neural network feature extractor as kernel hyperparameters, known as **deep kernel learning (DKL)**.




Outline

- 1 Introduction
- 2 Backgrounds
- 3 Methods
- 4 Results**
- 5 Potential Extensions
- 6 Conclusion


Overview




(a) Pendulum




(b) 2-DOF Reacher



(c) 7-DOF Pusher





(d) Swimmer



(e) Half Cheetah



Successful Cases

Certain variants of GP-TD-MPC outperformed the baseline:



Failed Cases

For the Swimmer task, only GP-TD-MPC with the Matérn kernel and DKL delivered comparable performance in the end.



Key Findings



From the experimental results, we derive the following observations:

- ① GP-TD-MPC with DKL typically requires fewer inducing points to achieve the error tolerance for the pivoted Cholesky method.
- ② GP-TD-MPC with the Matérn-3/2 kernel consistently matches or even outperforms the TD-MPC baseline.

Number of Inducing Points

The **pivoted Cholesky method** selects up to m inducing points and may terminate early if the truncation error falls below the specified tolerance.

For $m \leq 256$, only DKL variants reached the default tolerance.




Modeling Contact Dynamics

The spectral densities of the RBF kernel and the Matérn kernel are

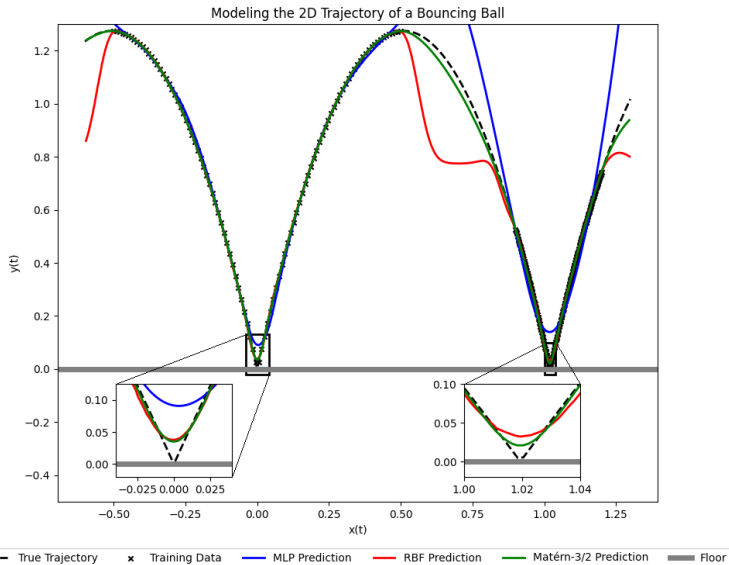
$$p_l^{\text{RBF}}(s) = (2\pi l^2)^{D/2} \exp(-2\pi^2 l^2 s^2)$$

$$p_{l,\nu}^{\text{Matérn}}(s) = \frac{2^D \pi^{D/2} \Gamma(\nu + D/2) (2\nu)^\nu}{\Gamma(\nu) l^{2\nu}} \left(\frac{2\nu}{l^2} + 4\pi^2 s^2 \right)^{-(\nu + D/2)}$$

resembling the Gaussian distribution and the t -distribution.



Toy Example: Trajectory of a Bouncing Ball



Runtime and Efficiency

Task Name	TD-MPC	RBF		Matérn-3/2		Spectral Mixture	
		Standard	DKL	Standard	DKL	Standard	DKL
Pendulum	175.13	260.31	266.87	279.64	297.81	430.92	287.64
Reacher	127.96	172.53	173.18	173.31	174.48	856.31	182.29
Pusher	201.40	305.83	293.66	315.20	302.04	—	—
Swimmer ($m \leq 256$)	1049.06	1531.21	1603.92	1537.10	1635.02	—	—
Swimmer ($m \leq 1024$)	1049.06	1592.99	1603.02	1666.68	1732.88	—	—
Half Cheetah	1064.46	2019.56	1631.59	2794.53	2818.31	—	—

Table 1: Comparison of total runtime (in seconds) averaged across 5 trials.

Outline

- 1 Introduction
- 2 Backgrounds
- 3 Methods
- 4 Results
- 5 Potential Extensions**
- 6 Conclusion

Motivation: For high-dimensional tasks that require a larger amount of data, more inducing points might be required. However, this may slow down inference. Instead, we consider alternative inference methods.



- ① **Local Kernel Interpolation.** $O(1)$ -time inference, scaling more effectively with a larger number of inducing points;
- ② **Dynamical Local Projection.** Making more efficient use of the limited inducing points.

Local Kernel Interpolation

Use **cubic interpolation** to obtain sparse matrix W s.t. $WK_{XX} \approx K_{XZ}$ and hence $K_{XX} \approx K_{XZ}K_{ZZ}^{-1}K_{ZX} \approx W^\top K_{ZZ}W$.



At any test location \mathbf{x}^* , compute the sparse interpolation vector $\mathbf{w}_{\mathbf{x}^*}$ s.t. $K_{ZZ}\mathbf{w}_{\mathbf{x}^*} \approx \mathbf{k}_{Z\mathbf{x}^*}$, we can approximate the original GP correction by

$$\text{output}(\mathbf{x}^*) = f_\theta(\mathbf{x}^*) + \underbrace{\mathbf{w}_{\mathbf{x}^*}^\top K_{ZZ}W(W^\top K_{ZZ}W + \sigma_\epsilon^2 \mathbf{I})^{-1}(\mathbf{y} - f_\theta(X))}_{\text{LKI Correction}}.$$



Dynamical Local Projection

Inspired by **online variational conditioning (OVC)**⁴ and **dynamic sparse GPs for MPC**⁵.



⁴Wesley J. Maddox, Samuel Stanton, and Andrew Gordon Wilson. *Conditioning Sparse Variational Gaussian Processes for Online Decision-making*. 2021.

⁵Lukas Hewing, Juraj Kabzan, and Melanie N. Zeilinger. "Cautious Model Predictive Control Using Gaussian Process Regression". In: (2020).

Dynamical Local Projection

Given a set of M inducing points Z , we compute the corresponding \mathbf{c} and C , they can thus be projected to a subset $Z' \subseteq Z$ of size m by

$$\begin{aligned}\mathbf{c}' &= K_{Z'X} \Sigma_{\mathbf{y}}^{-1} \mathbf{y} \approx K_{Z'Z} (K_{ZZ}^{-1} \mathbf{c}), \\ C' &= K_{Z'X} \Sigma_{\mathbf{y}}^{-1} K_{XZ'} \approx K_{Z'Z} (K_{ZZ}^{-1} C K_{ZZ}^{-1}) K_{ZZ'}.\end{aligned}$$

to compute the optimal variational distribution w.r.t. Z' . Each time, DLP takes $O(M^2 m)$ time, independent of the dataset size.

Dynamical Local Projection

Given a set of M inducing points Z , we compute the corresponding \mathbf{c} and C , they can thus be projected to a subset $Z' \subseteq Z$ of size m by

$$\begin{aligned}\mathbf{c}' &= K_{Z'X} \Sigma_{\mathbf{y}}^{-1} \mathbf{y} \approx K_{Z'Z} (K_{ZZ}^{-1} \mathbf{c}), \\ C' &= K_{Z'X} \Sigma_{\mathbf{y}}^{-1} K_{XZ'} \approx K_{Z'Z} (K_{ZZ}^{-1} C K_{ZZ}^{-1}) K_{ZZ'}.\end{aligned}$$


to compute the optimal variational distribution w.r.t. Z' . Each time, DLP takes $O(M^2 m)$ time, independent of the dataset size.

Let \mathbf{p} be the reference path, we can select Z' according to


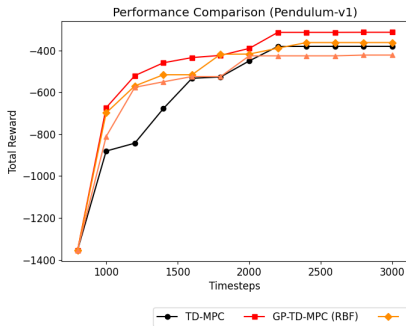
$$\text{cost}(\mathbf{z}, \mathbf{p}) = \min_{h=1, \dots, H-1} \eta^h \text{dist}(\mathbf{z}, \overline{\mathbf{p}_h \mathbf{p}_{h+1}})$$

which encourages more points to be selected near line segments corresponding to larger timesteps.

Dynamical Local Projection




Dynamical Local Projection



Future Work

We recommend some directions for future explorations:

- ① More informative **inducing point allocation (IPA)**.⁶
- ② Uncertainty quantification using **pathwise conditioning**.⁷
- ③ **Kernel composition** for domain-specific applications.



⁶Henry B. Moss, Sebastian W. Ober, and Victor Picheny. *Inducing Point Allocation for Sparse Gaussian Processes in High-Throughput Bayesian Optimisation*. 2023.

⁷James T. Wilson et al. *Pathwise Conditioning of Gaussian Processes*. 2021.

Outline

- 1 Introduction
- 2 Backgrounds
- 3 Methods
- 4 Results
- 5 Potential Extensions
- 6 Conclusion**